

实验报告

孔静 2014K8009929022

November 9, 2017

1 实验题目

路由器转发实验

2 实验内容

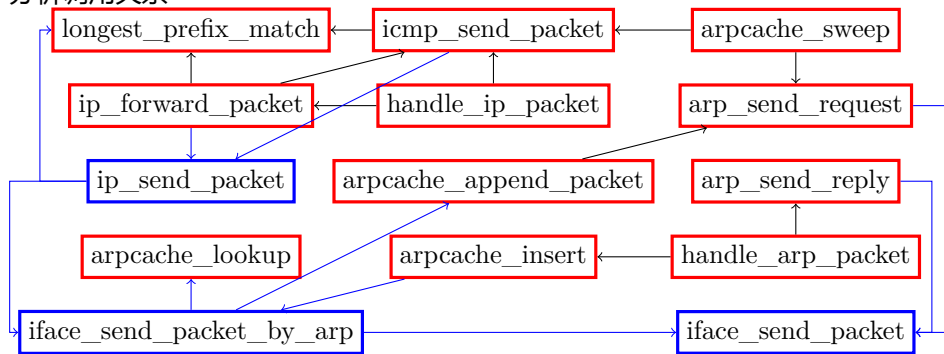
- 运行给定网络拓扑 (topo/router_topo.py)
- 在 r1 上执行路由器程序
 - 首先, 执行脚本 (disable_arp.sh, disable_icmp.sh, disable_ip_forward.sh), 禁止协议栈的相应功能
 - 在 r1 中运行 ./router, 进行数据包的处理
- 在 h1 上进行 ping 实验
 - Ping 10.0.1.1 (r1), 能够 ping 通
 - Ping 10.0.2.22 (h2), 能够 ping 通
 - Ping 10.0.3.33 (h3), 能够 ping 通
 - Ping 10.0.2.11, 返回 ICMP Destination Host Unreachable
 - Ping 10.0.4.1, 返回 ICMP Destination Net Unreachable

3 实验流程

- 确认函数功能
 - longest_prefix_match
 - ip_forward_packet
 - handle_ip_packet
 - icmp_send_packet

- arpcache_lookup
- arpcache_append_packet
- arpcache_insert
- arpcache_sweep
- arp_send_request
- arp_send_reply
- handle_arp_packet

• 分析调用关系



• 编写函数代码

```

1  //////////////////////////////////////////////////
2  //ip.c//
3  //////////////////////////////////////////////////
4  rt_entry_t *longest_prefix_match(u32 dst)
5  {
6      // fprintf(stderr, "TODO: longest prefix match for the packet.\n
7      ");
8      // return NULL;
9      rt_entry_t *entry, *longest = NULL;
10     list_for_each_entry(entry, &rttable, list){
11         if((dst & entry->mask) == (entry->dest & entry->mask)){
12             if(longest == NULL){
13                 longest = entry;
14             }
15             else if(entry->mask > longest->mask){
16                 longest = entry;
17             }
18         }
19     }
20     return longest;
21 }
22 void ip_forward_packet(u32 ip_dst, char *packet, int len)
23 {
24     // fprintf(stderr, "TODO: forward ip packet.\n");
25     struct iphdr *ip = packet_to_ip_hdr(packet);
26     rt_entry_t *entry = longest_prefix_match(ip_dst);
27     if(!entry){
28         icmp_send_packet(packet, len, (u8)3, (u8)0);
29     }
30 }

```

```

28     free(packet);
29     return;
30 }
31 ip->ttl--;
32 ip->checksum = ip_checksum(ip);
33 if(ip->ttl <= 0){
34     icmp_send_packet(packet, len, (u8)11, (u8)0);
35     free(packet);
36     return;
37 }
38 ip_send_packet(packet, len);
39 }
40 void handle_ip_packet(iface_info_t *iface, char *packet, int len)
41 {
42     // fprintf(stderr, "TODO: handle ip packet: echo the ping packet
43     // , and forward other IP packets.\n");
44     struct iphdr *ip = packet_to_ip_hdr(packet);
45     u32 ip_dst = ntohl(ip->daddr);
46     if(iface->ip == ip_dst){
47         icmp_send_packet(packet, len, (u8)0, (u8)0);
48         free(packet);
49     }
50     else{
51         ip_forward_packet(ip_dst, packet, len);
52     }
53 }
54 //icmp.c//
55 //icmp.c//
56 void icmp_send_packet(const char *in_pkt, int len, u8 type, u8
57     code)
58 {
59     // fprintf(stderr, "TODO: malloc and send icmp packet.\n");
60     struct iphdr *in_ip = packet_to_ip_hdr(in_pkt);
61     u32 dst = ntohl(in_ip->saddr);
62     rt_entry_t *entry = longest_prefix_match(dst);
63     int out_len, source, destination, new_len;
64     switch(type){
65     case 0:
66         out_len = len - IP_HDR_SIZE(in_ip) + IP_BASE_HDR_SIZE;
67         source = 0; destination = IP_HDR_SIZE(in_ip);
68         break;
69     case 3:
70     case 11:
71         out_len = ETHER_HDR_SIZE + IP_BASE_HDR_SIZE + ICMP_HDR_SIZE
72         + IP_HDR_SIZE(in_ip) + 8;
73         source = ICMP_HDR_SIZE; destination = 0;
74         break;
75     }
76     char *out_pkt;
77     out_pkt = malloc(out_len);
78     struct icmphdr* icmp_hdr = (struct icmphdr *) (out_pkt +
79     ETHER_HDR_SIZE + IP_BASE_HDR_SIZE);
80     memset(icmp_hdr, 0, ICMP_HDR_SIZE);
81     new_len = out_len - IP_BASE_HDR_SIZE - ETHER_HDR_SIZE;
82     memcpy((char *)icmp_hdr + source, (char *)in_ip + destination,
83     new_len - source);

```

```

80 icmp_hdr->type = type;
81 icmp_hdr->code = code;
82 icmp_hdr->checksum = icmp_checksum(icmp_hdr, new_len);
83 struct ether_header *out_eh = (struct ether_header *)out_pkt;
84 out_eh->ether_type = ntohs(ETH_P_IP);
85 struct iphdr *out_ip = packet_to_ip_hdr(out_pkt);
86 ip_init_hdr(out_ip, entry->iface->ip, dst, out_len -
    ETHER_HDR_SIZE, 1);
87 ip_send_packet(out_pkt, out_len);
88 }
89 ///////////////
90 //arp.c//
91 ///////////////
92 //request和reply功能相似, 顾合并成一个函数
93 void arp_send_option(iface_info_t *iface, u32 dst_ip, struct
    ether_arp *req_hdr, u16 option){
94
95     char *packet;
96     int len = ETHER_HDR_SIZE + sizeof(struct ether_arp);
97     packet = (char *)malloc(len);
98     struct ether_header* eh = (struct ether_header*)packet;
99     struct ether_arp* eh_arp = (struct ether_arp*)(packet +
    ETHER_HDR_SIZE);
100    eh->ether_type = htons(ETH_P_ARP);
101    memcpy((char *)eh->ether_shost, (char *)iface->mac, ETH_ALEN);
102    eh_arp->arp_hrd = htons(0x01);
103    eh_arp->arp_pro = htons(ETH_P_IP);
104    eh_arp->arp_hln = ETH_ALEN;
105    eh_arp->arp_pln = 4;
106    eh_arp->arp_op = htons(option);
107    eh_arp->arp_spa = htonl(iface->ip);
108
109    memset(eh_arp->arp_tha, 0, ETH_ALEN);
110    memcpy((char *)eh_arp->arp_sha, (char *)iface->mac, ETH_ALEN);
111    switch(option){
112        case ARPOP_REQUEST:
113            memset(eh->ether_dhost, 0xff, ETH_ALEN);
114            eh_arp->arp_tpa = htonl(dst_ip);
115            break;
116        case ARPOP_REPLY:
117            memcpy(eh->ether_dhost, req_hdr->arp_sha, ETH_ALEN);
118            eh_arp->arp_tpa = req_hdr->arp_spa;
119            break;
120    }
121    iface_send_packet(iface, packet, len);
122 }
123 void arp_send_request(iface_info_t *iface, u32 dst_ip)
124 {
125     // fprintf(stderr, "TODO: send arp request when lookup failed in
    arpcache.\n");
126     arp_send_option(iface, dst_ip, NULL, ARPOP_REQUEST);
127 }
128 void arp_send_reply(iface_info_t *iface, struct ether_arp *req_hdr
    )
129 {
130     // fprintf(stderr, "TODO: send arp reply when receiving arp
    request.\n");

```

```

131     arp_send_option(iface, 0, req_hdr, ARPOP_REPLY);
132 }
133 void handle_arp_packet(iface_info_t *iface, char *packet, int len)
134 {
135     // fprintf(stderr, "TODO: process arp packet: arp request & arp
136     // reply.\n");
137     struct ether_arp* eh_arp = (struct ether_arp*)(packet +
138     ETHER_HDR_SIZE);
139     u16 op = ntohs(eh_arp->arp_op);
140     u32 ip = ntohl(eh_arp->arp_tpa);
141
142     if(op == ARPOP_REQUEST && ip == iface->ip){
143         arp_send_reply(iface, eh_arp);
144     }
145     else if(op == ARPOP_REPLY && ip == iface->ip){
146         arpcache_insert(ntohl(eh_arp->arp_spa), eh_arp->arp_sha);
147     }
148 }
149 ///////////////////////////////////////////////////
150 //arpcache.c//
151 ///////////////////////////////////////////////////
152 int arpcache_lookup(u32 ip4, u8 mac[ETH_ALEN])
153 {
154     // fprintf(stderr, "TODO: lookup ip address in arp cache.\n");
155     // return 0;
156     int result = 0;
157     pthread_mutex_lock(&arpcache.lock);
158     for(int i = 0; i < MAX_ARP_SIZE; i++){
159         if(ip4 == arpcache.entries[i].ip4 && arpcache.entries[i].
160         valid){
161             memcpy(mac, arpcache.entries[i].mac, ETH_ALEN);
162             result = 1;
163             break;
164         }
165     }
166     pthread_mutex_unlock(&arpcache.lock);
167     return result;
168 }
169 void arpcache_append_packet(iface_info_t *iface, u32 ip4, char *
170 packet, int len)
171 {
172     // fprintf(stderr, "TODO: append the ip address if lookup failed
173     // , and send arp request if necessary.\n");
174     pthread_mutex_lock(&arpcache.lock);
175     struct arp_req *req_entry;
176     struct cached_pkt *pkt_entry;
177     list_for_each_entry(req_entry, &arpcache.req_list, list){
178         if (ip4 == req_entry->ip4){
179             goto append_packet;
180         }
181     }
182     req_entry = (struct arp_req *)malloc(sizeof(struct arp_req));
183     init_list_head(&req_entry->list);
184     list_add_tail(&req_entry->list, &arpcache.req_list);
185     init_list_head(&req_entry->cached_packets);
186     req_entry->iface = (iface_info_t *)malloc(sizeof(iface_info_t)
187 );

```

```

182     memcpy(req_entry->iface, iface, sizeof(iface_info_t));
183     req_entry->ip4 = ip4;
184     req_entry->sent = 0;
185     req_entry->retries = 0;
186     arp_send_request(iface, ip4);
187 append_packet:
188     pkt_entry = (struct cached_pkt *)malloc(sizeof(struct
189         cached_pkt));
189     init_list_head(&pkt_entry->list);
190     list_add_tail(&pkt_entry->list, &req_entry->cached_packets);
191     pkt_entry->packet = packet;
192     // pkt_entry->packet = (char *)malloc(len);
193     // memcpy(pkt_entry->packet, packet, len);
194     pkt_entry->len = len;
195
196     pthread_mutex_unlock(&arpcache.lock);
197 }
198 void arpcache_insert(u32 ip4, u8 mac[ETH_ALEN])
199 {
200     // fprintf(stderr, "TODO: insert ip->mac entry, and send all the
201         pending packets.\n");
201     pthread_mutex_lock(&arpcache.lock);
202     struct arp_req *req_entry, *req_q;
203     list_for_each_entry_safe(req_entry, req_q, &(arpcache.req_list),
204         list) {
204         if(req_entry->ip4 == ip4){
205             break;
206         }
207     }
208
209     int index;
210     for(index = 0; index < MAX_ARP_SIZE; index++){
211         if(arpcache.entries[index].valid == 0){
212             break;
213         }
214     }
215     if(index == MAX_ARP_SIZE){
216         index = rand() % MAX_ARP_SIZE;
217     }
218     memcpy(arpcache.entries[index].mac, mac, ETH_ALEN);
219     arpcache.entries[index].ip4 = ip4;
220     arpcache.entries[index].valid = 1;
221     arpcache.entries[index].added = 0;
222     pthread_mutex_unlock(&arpcache.lock);
223     struct cached_pkt *pkt_entry, *pkt_q;
224     list_for_each_entry_safe(pkt_entry, pkt_q, &(req_entry->
225         cached_packets), list) {
225         iface_send_packet_by_arp(req_entry->iface, ip4, pkt_entry->
226             packet, pkt_entry->len);
226         free(pkt_entry);
227     }
228     list_delete_entry(&(req_entry->list));
229     free(req_entry->iface);
230     free(req_entry);
231 }
232 void *arpcache_sweep(void *arg)
233 {

```

```

234 struct arp_req *req_entry , *req_q;
235 struct cached_pkt *pkt_entry , *pkt_q;
236 while (1) {
237     sleep(1);
238     // fprintf(stderr, "TODO: sweep arpcache periodically: remove
239     // old entries, resend arp requests .\n");
240     pthread_mutex_lock(&arpcache.lock);
241     for(int index = 0; index < MAX_ARP_SIZE; index++){
242         if(++arpcache.entries[index].added > ARP_ENTRY_TIMEOUT){
243             arpcache.entries[index].valid = 0;
244         }
245     }
246     list_for_each_entry_safe(req_entry , req_q, &(arpcache.req_list
247     ), list) {
248         if(req_entry->retries >= ARP_REQUEST_MAX_RETRIES){
249             list_delete_entry(&(req_entry->list));
250             pthread_mutex_unlock(&arpcache.lock);
251             list_for_each_entry_safe(pkt_entry , pkt_q, &(req_entry->
252             cached_packets), list){
253                 icmp_send_packet(pkt_entry->packet , pkt_entry->len , (u8)
254                 3, (u8)1);
255                 free(pkt_entry->packet);
256                 free(pkt_entry);
257             }
258             pthread_mutex_lock(&arpcache.lock);
259             free(req_entry->iface);
260             free(req_entry);
261         }
262         else{
263             arp_send_request(req_entry->iface , req_entry->ip4);
264             req_entry->sent = time(NULL);
265             req_entry->retries += 1;
266         }
267     }
268     pthread_mutex_unlock(&arpcache.lock);
269 }

```

• 进行结果测试

```

1 #略去一些代码
2 os.system("make clean")
3 os.system("make")
4 net.start()
5 h1, r1 = net.get('h1', 'r1')
6
7 r1.cmd('./scripts/disable_arp.sh')
8 r1.cmd('./scripts/disable_icmp.sh')
9 r1.cmd('./scripts/disable_ip_forward.sh')
10 r1.cmd('./router &')
11 time.sleep(0.2)
12 print h1.cmd('ping 10.0.1.1 -c 5')
13 print h1.cmd('ping 10.0.2.22 -c 5')
14 print h1.cmd('ping 10.0.3.33 -c 5')

```

```

15 print h1.cmd('ping 10.0.2.11 -c 5')
16 print h1.cmd('ping 10.0.4.1 -c 5')
17 net.stop()

```

4 实验结果

```

1 kj@12-ubuntu:~/Desktop/07-router$ sudo python test.py
2 rm -f *.o router
3 gcc -c -g -Wall -Iinclude arp.c -o arp.o
4 In file included from arp.c:12:0:
5 include/log.h:12:23: warning: 'this_log_level' defined but not used
6 [-Wunused-variable]
7 static enum log_level this_log_level = DEBUG;
8 ^
9 include/log.h:14:20: warning: 'log_level_str' defined but not used [-
10 Wunused-variable]
11 static const char *log_level_str[] = { "DEBUG", "INFO", "WARNING", "
12 ERROR" };
13 ^
14 gcc -c -g -Wall -Iinclude arpcache.c -o arpcache.o
15 gcc -c -g -Wall -Iinclude icmp.c -o icmp.o
16 gcc -c -g -Wall -Iinclude ip.c -o ip.o
17 gcc -c -g -Wall -Iinclude main.c -o main.o
18 gcc -c -g -Wall -Iinclude packet.c -o packet.o
19 gcc -c -g -Wall -Iinclude rtable.c -o rtable.o
20 gcc arp.o arpcache.o icmp.o ip.o main.o packet.o rtable.o -o router -
21 lpthread
22 PING 10.0.1.1 (10.0.1.1) 56(84) bytes of data.
23 64 bytes from 10.0.1.1: icmp_seq=1 ttl=64 time=0.130 ms
24 64 bytes from 10.0.1.1: icmp_seq=2 ttl=64 time=0.051 ms
25 64 bytes from 10.0.1.1: icmp_seq=3 ttl=64 time=0.084 ms
26 64 bytes from 10.0.1.1: icmp_seq=4 ttl=64 time=0.233 ms
27 64 bytes from 10.0.1.1: icmp_seq=5 ttl=64 time=0.118 ms
28
29 — 10.0.1.1 ping statistics —
30 5 packets transmitted, 5 received, 0% packet loss, time 4103ms
31 rtt min/avg/max/mdev = 0.051/206.291/1030.972/412.340 ms, pipe 2
32
33 PING 10.0.2.22 (10.0.2.22) 56(84) bytes of data.
34 64 bytes from 10.0.2.22: icmp_seq=1 ttl=63 time=0.147 ms
35 64 bytes from 10.0.2.22: icmp_seq=2 ttl=63 time=0.057 ms
36 64 bytes from 10.0.2.22: icmp_seq=3 ttl=63 time=0.100 ms
37 64 bytes from 10.0.2.22: icmp_seq=4 ttl=63 time=0.155 ms
38 64 bytes from 10.0.2.22: icmp_seq=5 ttl=63 time=0.140 ms
39
40 — 10.0.2.22 ping statistics —
41 5 packets transmitted, 5 received, 0% packet loss, time 4092ms
42 rtt min/avg/max/mdev = 0.057/0.119/0.155/0.039 ms
43
44 PING 10.0.3.33 (10.0.3.33) 56(84) bytes of data.
45 64 bytes from 10.0.3.33: icmp_seq=1 ttl=63 time=0.077 ms
46 64 bytes from 10.0.3.33: icmp_seq=2 ttl=63 time=0.067 ms
47 64 bytes from 10.0.3.33: icmp_seq=3 ttl=63 time=0.088 ms
48 64 bytes from 10.0.3.33: icmp_seq=4 ttl=63 time=0.052 ms

```



```

45 64 bytes from 10.0.3.33: icmp_seq=5 ttl=63 time=0.140 ms
46
47 — 10.0.3.33 ping statistics —
48 5 packets transmitted, 5 received, 0% packet loss, time 4091ms
49 rtt min/avg/max/mdev = 0.052/0.084/0.140/0.032 ms
50
51 PING 10.0.2.11 (10.0.2.11) 56(84) bytes of data.
52 From 10.0.1.1 icmp_seq=1 Destination Host Unreachable
53
54 — 10.0.2.11 ping statistics —
55 5 packets transmitted, 0 received, +1 errors, 100% packet loss, time
    4087ms
56 pipe 5
57
58 PING 10.0.4.1 (10.0.4.1) 56(84) bytes of data.
59 From 10.0.1.1 icmp_seq=1 Destination Net Unreachable
60 From 10.0.1.1 icmp_seq=2 Destination Net Unreachable
61 From 10.0.1.1 icmp_seq=3 Destination Net Unreachable
62 From 10.0.1.1 icmp_seq=4 Destination Net Unreachable
63 From 10.0.1.1 icmp_seq=5 Destination Net Unreachable
64
65 — 10.0.4.1 ping statistics —
66 5 packets transmitted, 0 received, +5 errors, 100% packet loss, time
    4077ms

```

5 结果分析

ping 命令会通过 ICMP 协议发送 ICMP 数据包, 步骤如下:

1. 应用程序构造数据包, 该示例是产生 ICMP 包, 被提交给内核 (网络驱动程序)
2. 内核检查是否能够转化该 IP 地址为 MAC 地址, 也就是在本地的 ARP 缓存中查看 IP-MAC 对应表
3. 如果存在该 IP-MAC 对应关系, 那么跳到步骤 7; 否则继续以下步骤
4. 内核进行 ARP 广播, 目的地的 MAC 地址是 FF-FF-FF-FF-FF-FF, ARP 命令类型为 REQUEST(1), 其中包含有自己的 MAC 地址
5. 当 192.168.1.2 主机接收到该 ARP 请求后, 就发送一个 ARP 的 REPLY(2) 命令, 其中包含自己的 MAC 地址
6. 本地获得 192.168.1.2 主机的 IP-MAC 地址对应关系, 并保存到 ARP 缓存中
7. 内核将把 IP 转化为 MAC 地址, 然后封装在以太网头结构中, 再把数据发送出去